




Avoiding the Tower of Babel

*Effectively evolving events in
micro-service ecosystems*

Sebastián Ortega

letgo



Founded in 2015

Phenomenal growth

80 MM+ downloads

Focused in USA and Turkey

Offices in Barcelona and NY



SPLIT THE MONOLITH!



events at letgo by
Rafa Gómez

A μ -service is...

- it's own process
- small
- independently deployed
- networked
 - API calls
 - pub/sub of events


Are service architectures scalable?

- No intrinsic advantage: you can distribute your monolith
- Except for having more **freedom** to scale independently each service



Are service architectures scalable?

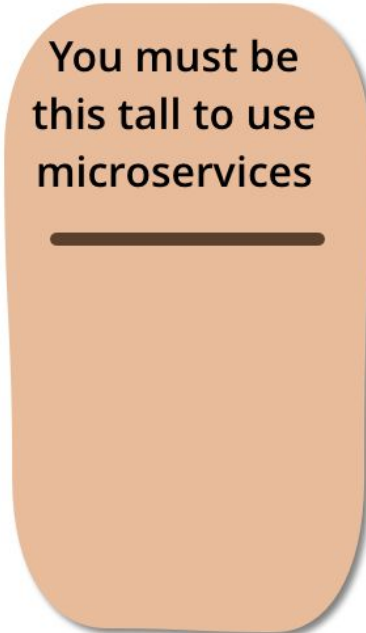
- **Social scalability**: the true scalability property of services
- Conway's law in reverse
- Squads: team with **all the roles** and empowered through **independence**
 - Prefer service **choreography** to service **orchestration**



MELVIN CONWAY

Is this for free?

- NOPE
- more complex to operate
- bigger footprint
- prone to (excessive) heterogeneity
- challenging to manage change at the boundaries



You must be
this tall to use
microservices



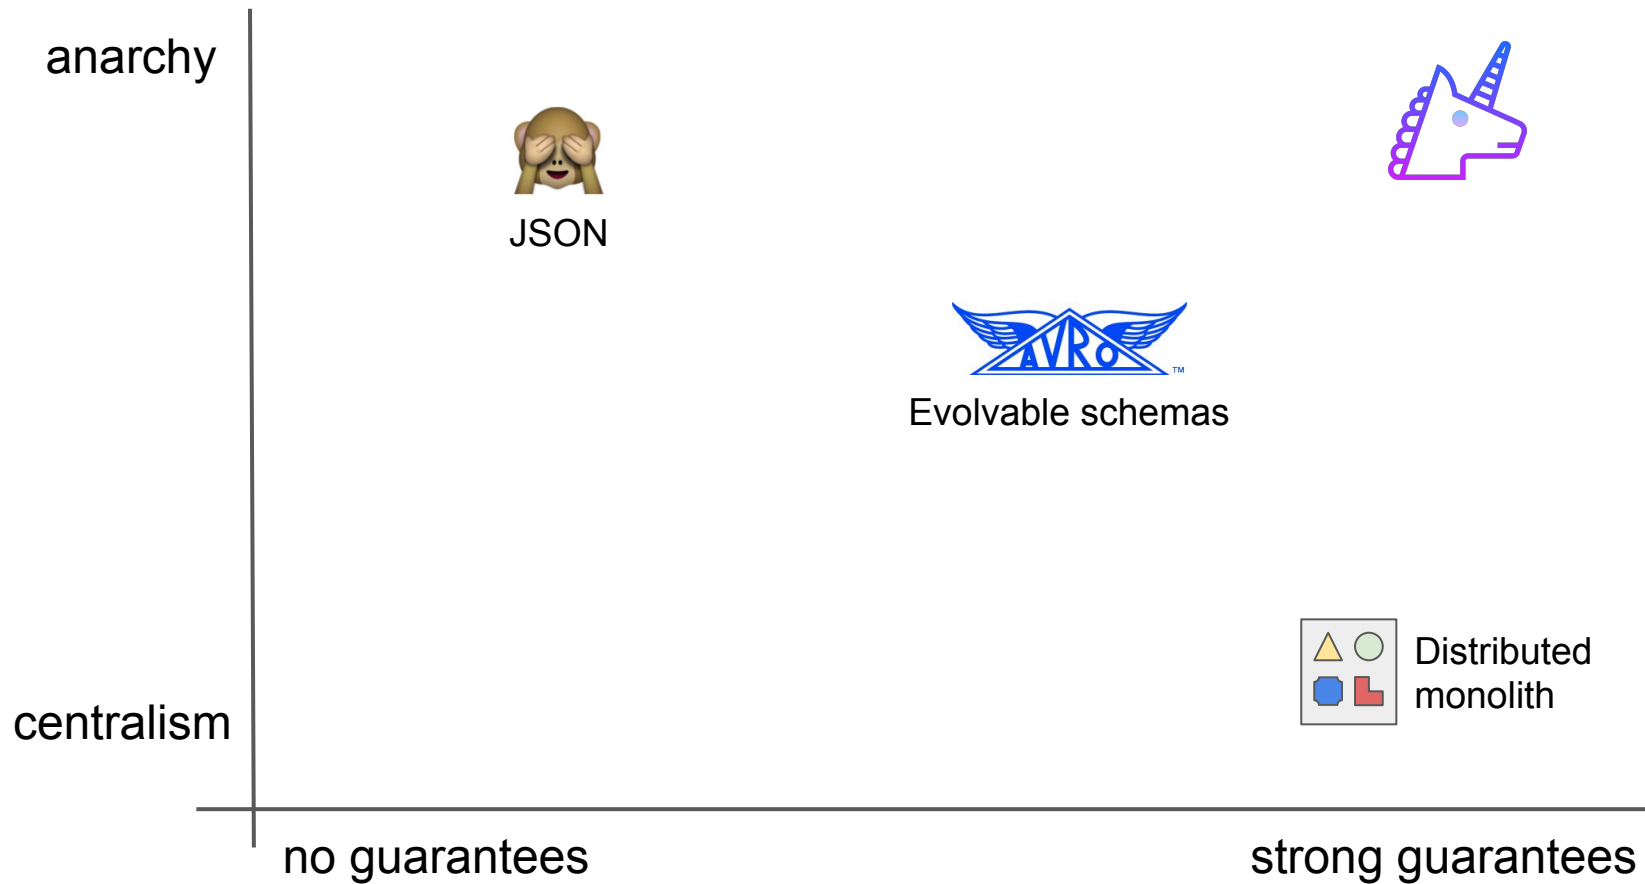


Don't even consider microservices
unless you have a system that's
too complex to manage as a
monolith.

—Martin Fowler



HOW TO AVOID BABEL WITHOUT CURTAILING INDEPENDENCE?



Apache AVRO



- Serialization protocol
- Schema-based
- Binary
 - Fast parsing!
 - Compact
 - Tools to inspect :(
- Polyglot: java, scala, js, python, C#...
- Evolvable schemas!

```
{  
  "type": "record",  
  "namespace": "com.tidx",  
  "name": "ProductVisit",  
  "fields": [  
    {  
      "name": "product_id",  
      "type": "string",  
      "doc": "Identifier of  
the visited product"  
    },  
    ...  
  ]  
}
```

Evolvable schemas

- Not every schema will be easy to evolve
- Keep required fields to the minimum
- Provide default values when possible
- Introduce only optional/with defaults fields
- Let's see an example

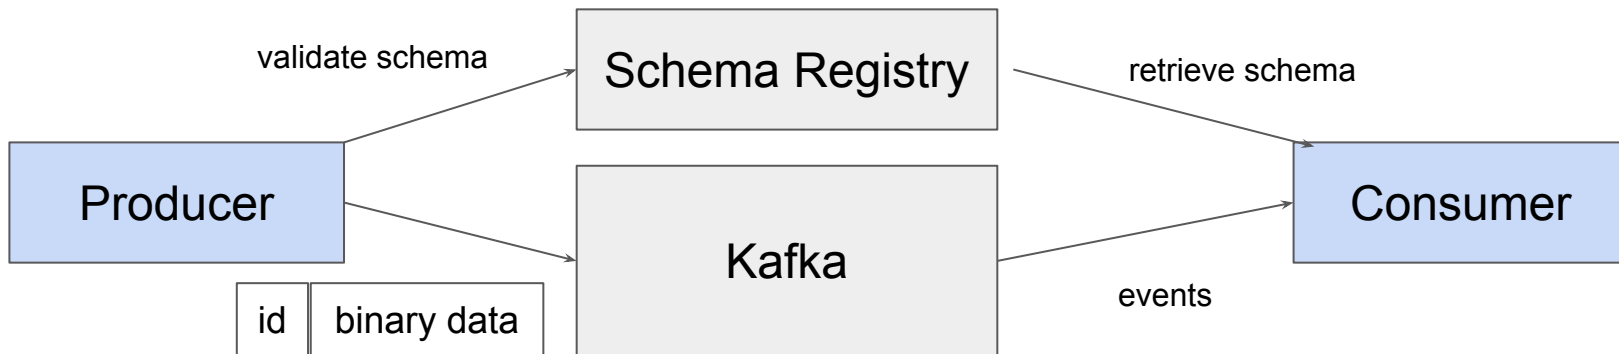


Apache Kafka

- Distributed commit log open sourced by LinkedIn
- Its creators founded Confluent to provide commercial support
- Kafka offers
 - exactly-once delivery of events
 - high-throughput (millions of events per second)
 - high-reliability
- How?
 - for kafka, payloads are just byte arrays
 - no parsing, zero-copy optimizations
- So, Kafka is blind regarding schemas and schema validation...

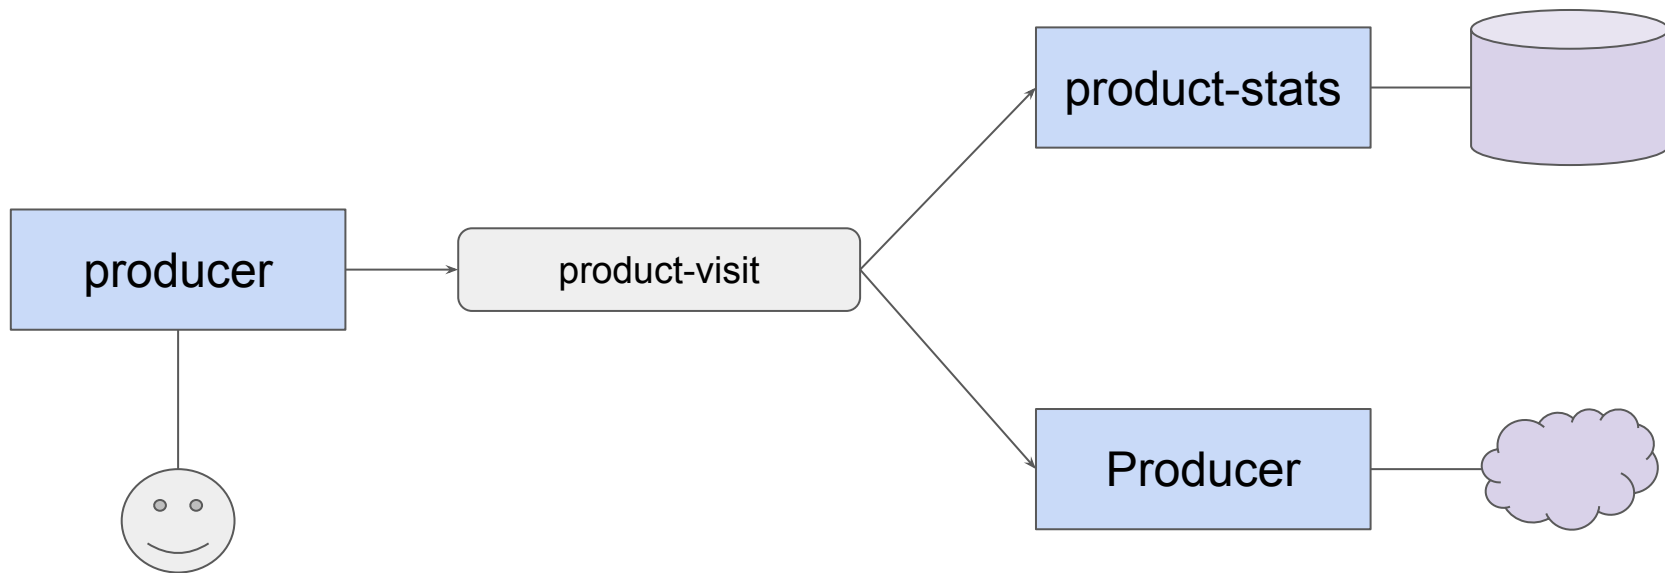
Kafka - Schema Registry

- Component developed by Confluent (also OSS)
- Used at the edges of Kafka
- Enforces [schema compatibility](#): no need to cut red tape



Demo time

[code at github](#)



GIVE AUTONOMY AND THE BEST
COLLABORATION TOOLS TO YOUR SQUADS
AND THEY'LL MOVE THE WORLD.





Thank you

@_sortega

(letgo is hiring)

we.letgo.com/contact